

Distributed Video Slow Motion using Scene-Based Frame Interpolation

Team Members

Anisha Priyadarshini Mohanty (apmohant@uci.edu)

Venkata Bharath Reddy Reddem (vreddem@uci.edu)

Maitreyi Sinha (maitres@uci.edu)



Background

Problem: Most videos are captured at 24-30 FPS, which is too low for smooth motion or slow-motion playback

One solution:

Frame interpolation generates new frames between existing ones to produce high-FPS, smoother video

- Modern deep learning models: RIFE create high-quality results but are **computationally expensive**

Videos have scenes with varying motion making some segments much costlier to process

- Static scenes are cheap to process
- Fast-motion scenes require more intermediate frames



This variability makes interpolation **slow and imbalanced** on a single machine or with naive equal splitting

Introduction

- Video processing is **computationally expensive**.
- High quality slow motion video requires frame interpolation.
- Interpolating entire video sequentially is time consuming for long videos.
- Running the entire pipeline on a single machine leads to slow processing times.
- Distributing work across multiple nodes can significantly improve throughput.
- Motion complexity varies across scenes, making load balancing across workers difficult.



Motivation

- Slow motion is widely used in sports, video editing, and visual effects.
- Scenes in a video vary widely in motion complexity.
- Motion intensity varies across each set of frames:
 - High motion → more frames needed
 - Low motion → fewer frames suffice
- Distributed processing reduces computation time.
- Goal
 - Develop a distributed system to generate slow-motion videos efficiently.
 - Adaptively interpolate different number of frames per scene based on motion intensity.

$$\text{motion_intensity} = \frac{1}{N-1} \sum_{i=1}^{N-1} \|F_{i+1} - F_i\|$$

Related Works

- **Distributed Video Processing**
 - Simple frame splitting across workers.
 - Ignores motion complexity → severe load imbalance.
- **Scene-Based Partitioning**
 - Splits video into scenes for semantic consistency.
 - Assumes equal cost per frame, inaccurate for interpolation tasks.

Ref: [Jeffrey](#), [Huang](#), [Nandra](#)

- **Load Balancing in Distributed Pipelines**
 - Traditional methods: static partitioning or round-robin
 - Often fail when task cost per frame varies significantly.
- **Video Frame Interpolation Systems**
 - Focuses on model accuracy/quality, not distributed scheduling.
 - Computational cost highly depends on motion intensity

Ref: [Jiang et al. \(2018\)](#), [Bao et al. \(2019\)](#), [Huang et al. \(2022\)](#)

System Architecture: AWS Distributed Pipeline



- **Scene Detection:** Split video into multiple scenes based on scene change.
- **Motion Analysis:** Compute motion intensity per scene.
- **Adaptive Frame Interpolation:** More frames for high-motion scenes; fewer frames for low-motion scenes
- **Distributed Execution:** Assign scenes to worker nodes based on estimated load and run interpolation in parallel.


System Architecture Overview

- Master node analyzes video and performs scene detection.
- Motion intensity computed per scene.
- Scenes assigned to worker nodes based on estimated load.
- Workers run interpolation and upload results.
- Master merges all processed scenes.



Key Features

- Adaptive interpolation per scene based on motion intensity.
- Parallel processing to reduce total computation time.
- Use RIFE interpolation model which uses IFNet for estimating intermediate flows.
- S3 integration for seamless distributed workflow.
- Dynamic load balancing ensures optimal utilization of all worker nodes.
- Modular pipeline design makes the system easily scalable to larger video datasets.





Understanding Our Processing Pipeline


- Scene Detection
 - Histogram based scene boundary detection.
 - Automatically finds changes in scene transitions.
 - Outputs timestamps where new scenes begin.
 - Reduces unnecessary computation on static segments.
- Motion-Adaptive Load Balancing
 - Compute motion intensity for each scene.
 - Estimate number of interpolated frames needed.
 - Greedy load assignment ensures better balance.
 - Workers process nearly equal workload.

Results


Performance Comparison of Load-Balancing Strategies



Worker Utilization — Equal Split




Worker Utilization — Our Scene-Aware Balancing



Workers are utilized evenly with scene-based load assignment

Master Node Overhead Breakdown (Excluding Worker Execution)



Scene detection and dynamic assignments add almost no overhead.
(< 2% of processing time)


Conclusion

- Designed and implemented a distributed system that splits tasks across multiple worker nodes for faster processing
- Demonstrated that parallel execution significantly improves throughput, with measurable gains even on a two-node setup
- Identified key system challenges, including load imbalance, synchronization overhead, and worker heterogeneity

Future Work

- Extend support for heterogeneous hardware (CPUs, GPUs, accelerators)
- Add auto scaling based on workload intensity
- Integrate detailed monitoring (metrics, logs, tracing) for deeper performance analysis
- Implement fault-tolerance features such as task reassignment, checkpointing, and worker health checks

Demo



Thank You! Questions?

References

- Wenbo Bao, Wei-Sheng Lai, Chao Ma, Xiaoyun Zhang, Zhiyong Gao, and Ming-Hsuan Yang. Depth-aware video frame interpolation. In Proceedings of the IEEE/CVF conference on computer vision and pattern recognition, pp. 3703–3712, 2019.
- Zhewei Huang, Tianyuan Zhang, Wen Heng, Boxin Shi, and Shuchang Zhou. Real-time intermediate flow estimation for video frame interpolation. In European Conference on Computer Vision, pp.624–642. Springer, 2022.
- Huaizu Jiang, Deqing Sun, Varun Jampani, Ming-Hsuan Yang, Erik Learned-Miller, and Jan Kautz. Super slomo: High quality estimation of multiple intermediate frames for video interpolation. In Proceedings of the IEEE conference on computer vision and pattern recognition, pp. 9000–9008, 2018.
- Jeffrey Dean and Sanjay Ghemawat. MapReduce: Simplified data processing on large clusters. In *Proceedings of the USENIX Symposium on Operating Systems Design and Implementation*, pp. 137–150, 2004.
- Qi Huang, Petchean Ang, Peter Knowles, Tomasz Nykiel, Iaroslav Tverdokhlib, Amit Yajurvedi, Paul Dapolito IV, Xifan Yan, Maxim Bykov, and Chuen Liang. SVE: Distributed video processing at Facebook scale. In Proceedings of the 26th ACM Symposium on Operating Systems Principles (SOSP), pp. 87–103, 2017.
- Nandra, Constantin, Victor Bacu, and Dorian Gorgan. Distributed, workflow-driven rendering of 3D object scenes on a big data processing platform. In Proceedings of the 22nd International Conference on System Theory, Control and Computing (ICSTCC), pp. 94–99, 2018.
- Aghasoleymani Najafabadi, Saeid. Performance Assessment of Load Balancing Methods in Cloud Computing: Analysis of Round Robin, Equally Spread, and Throttled Strategies Using Cloud Analyst. arXiv preprint arXiv:2507.11899, 2025. Ahmed Hassani, Mohamed Elgarib, Ahmed Selim, Sung-Ho Bae, Mohamed Hefeeda, and Wojciech Matusik. Large-scale, fast and accurate shot boundary detection through spatio-temporal convolutional neural networks. arXiv preprint arXiv:1705.03281, 2017.